

[Google Discussions](#)

Exam Professional Data Engineer All Questions

View all questions & answers for the Professional Data Engineer exam

[Go to Exam](#)

EXAM PROFESSIONAL DATA ENGINEER TOPIC 1 QUESTION 228 DISCUSSION

Actual exam question from Google's Professional Data Engineer

Question #: 228

Topic #: 1

[\[All Professional Data Engineer Questions\]](#)

You have a streaming pipeline that ingests data from Pub/Sub in production. You need to update this streaming pipeline with improved business logic. You need to ensure that the updated pipeline reprocesses the previous two days of delivered Pub/Sub messages. What should you do? (Choose two.)

- A. Use the Pub/Sub subscription clear-retry-policy flag
- B. Use Pub/Sub Snapshot capture two days before the deployment.
- C. Create a new Pub/Sub subscription two days before the deployment.
- D. Use the Pub/Sub subscription retain-acked-messages flag.
- E. Use Pub/Sub Seek with a timestamp.

[Show Suggested Answer](#)

by [e70ea9e](#) at Dec. 30, 2023, 9:54 a.m.

Comments

[Submit](#)

[tibuenoc](#) [Highly Voted](#) 1 year, 3 months ago

Selected Answer: D

DE

Another way to replay messages that have been acknowledged is to seek to a timestamp. To seek to a timestamp, you must first configure the subscription to retain acknowledged messages using retain-acked-messages. If retain-acked-messages is set, Pub/Sub retains acknowledged messages for 7 days.

You only need to do this step if you intend to seek to a timestamp, not to a snapshot.

<https://cloud.google.com/pubsub/docs/replay-message>

   upvoted 16 times

  **nadavw** 8 months, 1 week ago



this is correct as there are 2 options (timestamp and snapshot) and foreach there are 2 stages.

Snapshot - create ('B') and seek

Timestamp - configure 'retain' ('D') and seek ('E')

as shown 'B' is missing the 'seek' operation

   upvoted 1 times

  **joao_01** 1 year ago

Its BE.

By the way, you can seek to a snapshot yes:

"Seeks an existing subscription to a point in time or to a given snapshot, whichever is provided in the request"

Link:<https://cloud.google.com/pubsub/docs/reference/rest/v1/projects.subscriptions/seek>

   upvoted 2 times

  **2ad2bc7** 4 months, 1 week ago




Think about it -

what you want to do - process last 2 days of messages.

What does snapshot give you - it give you what were the un-acknowledged messages in pub/sub at that point in time 2 days ago.

How will that help you process messages that were sent to pub/sub in the last 2 days (i.e. after the snapshot?)

   upvoted 2 times

  **GCP001** **Highly Voted**  1 year, 3 months ago

Selected Answer: B

B and E, already tested at cloud console.

   upvoted 12 times

  **desertlotus1211** **Most Recent**  1 month, 1 week ago

Selected Answer: B

The only way to ensure that the updated pipeline reprocesses the previous two days of delivered Pub/Sub messages (current) is to use B&E.... for future D&E can work.

   upvoted 1 times

  **71083a7** 3 months, 1 week ago

Selected Answer: D

Since E is seek to a timestamp only one to replay, I think D is correct answer and not B

   upvoted 1 times

  **AWSandeep** 4 months, 2 weeks ago

D and E


Please seem to not understand that the retain_acked_messages parameter needs to be enabled for any snapshot or seeking functionality to work.

   upvoted 1 times

  **SamuelTsch** 6 months, 1 week ago

Selected Answer: B



After reading https://cloud.google.com/pubsub/docs/replay-overview#snapshot_overview carefully, B and E are correct.

   upvoted 1 times

  **SamuelTsch** 6 months, 1 week ago

change my mind. D and E should be correct.

   upvoted 2 times

  **Preetmehta1234** 7 months, 1 week ago

Selected Answer: D

First, read this document: <https://cloud.google.com/pubsub/docs/replay-overview>.

Key Points:

Seek to a Snapshot: Reprocesses only unacknowledged messages.

Seek to a Timestamp: Reprocesses all messages (acknowledged and unacknowledged) after that time.

Since the question asks for delivering all messages, option E is correct, as it includes both acknowledged and unacknowledged messages.

Regarding Option D: Configuring a subscription with the `retain_acked_messages` property allows replaying previously acknowledged messages retained for up to 31 days. This satisfies the requirement to deliver all messages and retains them longer than the mentioned 2 days.

   upvoted 1 times

  **MithunDesai** 8 months, 3 weeks ago

Selected Answer: B

B&D - based on Vertex AI feedback

   upvoted 1 times

  **Anudeep58** 11 months ago

Selected Answer: E

BE

B. Use Pub/Sub Snapshot capture two days before the deployment.

Pub/Sub Snapshot: Creating a snapshot captures the state of the subscription at a specific point in time. You can then seek to this snapshot to replay messages from that point onwards.

By capturing a snapshot two days before the deployment, you can ensure that your pipeline reprocesses messages from the past two days.

E. Use Pub/Sub Seek with a timestamp.

Pub/Sub Seek: This feature allows you to reset the subscription to a specific timestamp. Messages published to the topic after this timestamp are re-delivered.

By seeking to the timestamp from two days ago, you can instruct Pub/Sub to start re-delivering messages from that point in time

   upvoted 2 times

  **virat_kohli** 11 months, 2 weeks ago

Selected Answer: D

D. Use the Pub/Sub subscription `retain-acked-messages` flag.

E. Use Pub/Sub Seek with a timestamp.

   upvoted 2 times

  **cuadradobertolinisebastiancami** 1 year, 2 months ago

Selected Answer: D

E for sure, you need to seek from a timestamp.

To accomplish that you need to "Set the `retain-acked-messages` flag to true for the subscription."

From google documentation:

"Note: To seek to a previous time point, your subscription must be configured to retain acknowledged messages. You can change this setting by clicking Edit on the subscription details page, and checking the box for Retain acknowledged messages."

<https://cloud.google.com/pubsub/docs/replay-message>

   upvoted 3 times

  **Tryolabs** 1 year, 2 months ago

Selected Answer: E

D and E,

<https://cloud.google.com/pubsub/docs/replay-message>

   upvoted 3 times

  **ML6** 1 year, 2 months ago

Selected Answer: E

B and E: The seek feature extends subscriber capabilities by allowing you to alter the acknowledgement state of messages in bulk. For example, you can replay previously acknowledged messages or purge messages in bulk. In addition, you can copy the acknowledgement state of one subscription to another by using seek in combination with a snapshot. Source:

<https://cloud.google.com/pubsub/docs/replay-overview>

👍 ↩ 🚩 upvoted 3 times

🗨️ 👤 **Sofia98** 1 year, 3 months ago

Selected Answer: B

BE

<https://cloud.google.com/pubsub/docs/replay-overview>

👍 ↩ 🚩 upvoted 3 times

🗨️ 👤 **tibuenoc** 1 year, 3 months ago

But There is a problem snapshot you should seek by subscriptions not by timestamp

👍 ↩ 🚩 upvoted 1 times

🗨️ 👤 **Matt_108** 1 year, 3 months ago

Selected Answer: D

Option D and E

👍 ↩ 🚩 upvoted 3 times

🗨️ 👤 **task_7** 1 year, 3 months ago

Selected Answer: D

DE

Set the retain-acked-messages flag to true for the subscription.

This instructs Pub/Sub to store acknowledged messages for a specified retention period.

E Use Pub/Sub Seek with a timestamp.

After deploying the updated pipeline, use the Seek feature to replay messages.

Specify a timestamp that's two days before the current time.

This rewinds the subscription's message cursor, making it redeliver messages from that point onward.

👍 ↩ 🚩 upvoted 6 times

🗨️ 👤 **raaad** 1 year, 4 months ago

Selected Answer: B

- Pub/Sub Snapshots allow you to capture the state of a subscription's unacknowledged messages at a particular point in time.

- By creating a snapshot two days before deploying the updated pipeline, you can later use this snapshot to replay the messages from that point.

=====

Option E:

- Pub/Sub Seek allows us to alter the acknowledgment state of messages in bulk.

- So we can rewind a subscription to a point in time or a snapshot.

- Using Seek with a timestamp corresponding to two days ago would allow the updated pipeline to reprocess messages from that time.

👍 ↩ 🚩 upvoted 7 times

🗨️ 👤 **datapassionate** 1 year, 3 months ago

Creating a snapshot of the Pub/Sub subscription two days before the deployment captures the state of unacknowledged messages at that particular point in time, which would include messages from before those two days. If our objective is to reprocess the data from the last two days specifically, then capturing a snapshot two days prior wouldn't directly address this need.

👍 ↩ 🚩 upvoted 3 times

🗨️ 👤 **datapassionate** 1 year, 3 months ago

This case is described here.

<https://cloud.google.com/pubsub/docs/replay-message>

And according to this D & E would be correct.

👍 ↩ 🚩 upvoted 5 times

🗨️ 👤 **datapassionate** 1 year, 3 months ago

Another way to replay messages that have been acknowledged is to seek to a timestamp. To seek to a timestamp, you must first configure the subscription to retain acknowledged messages using retain-acked-messages. If retain-acked-messages is set, Pub/Sub retains acknowledged messages for 7 days.

👍 ↩ 🚩 upvoted 4 times

[Load full discussion...](#)



Platform

> Home

> Examtopics PRO

> All Exams

> Training Courses



© 2024 ExamTopics