

[Google Discussions](#)

Exam Professional Data Engineer All Questions

View all questions & answers for the Professional Data Engineer exam

[Go to Exam](#)

EXAM PROFESSIONAL DATA ENGINEER TOPIC 1 QUESTION 267 DISCUSSION

Actual exam question from Google's Professional Data Engineer

Question #: 267

Topic #: 1

[\[All Professional Data Engineer Questions\]](#)

You are creating a data model in BigQuery that will hold retail transaction data. Your two largest tables, `sales_transaction_header` and `sales_transaction_line`, have a tightly coupled immutable relationship. These tables are rarely modified after load and are frequently joined when queried. You need to model the `sales_transaction_header` and `sales_transaction_line` tables to improve the performance of data analytics queries. What should you do?

- A. Create a `sales_transaction` table that holds the `sales_transaction_header` information as rows and the `sales_transaction_line` rows as nested and repeated fields.
- B. Create a `sales_transaction` table that holds the `sales_transaction_header` and `sales_transaction_line` information as rows, duplicating the `sales_transaction_header` data for each line.
- C. Create a `sales_transaction` table that stores the `sales_transaction_header` and `sales_transaction_line` data as a JSON data type.
- D. Create separate `sales_transaction_header` and `sales_transaction_line` tables and, when querying, specify the `sales_transaction_line` first in the WHERE clause.

[Show Suggested Answer](#)

by [scaenruy](#) at Jan. 3, 2024, 6:27 p.m.

Comments

Submit

🗂️ 👤 **raaad** Highly Voted 👍 10 months ago

Selected Answer: A

- In BigQuery, nested and repeated fields can significantly improve performance for certain types of queries, especially joins, because the data is co-located and can be read efficiently. - - This approach is often used in data warehousing scenarios where query performance is a priority, and the data relationships are immutable and rarely modified.

👍 ↩️ 🚩 upvoted 6 times

🗂️ 👤 **josech** Most Recent 🕒 5 months, 2 weeks ago

Selected Answer: A

Option A <https://cloud.google.com/bigquery/docs/best-practices-performance-nested>

👍 ↩️ 🚩 upvoted 2 times

🗂️ 👤 **hanoverquay** 7 months, 3 weeks ago

Selected Answer: A

optional A

👍 ↩️ 🚩 upvoted 1 times

🗂️ 👤 **JyoGCP** 8 months, 2 weeks ago

Selected Answer: A

Option A

👍 ↩️ 🚩 upvoted 1 times

🗂️ 👤 **Matt_108** 9 months, 3 weeks ago

Selected Answer: A

Option A

👍 ↩️ 🚩 upvoted 1 times

🗂️ 👤 **scaenruy** 10 months ago

Selected Answer: A

A. Create a sales_transaction table that holds the sales_transaction_header information as rows and the sales_transaction_line rows as nested and repeated fields.

👍 ↩️ 🚩 upvoted 1 times



Platform

> Home

> Examtopics PRO

> All Exams

> Training Courses

